

Power Consumption

Overview

In this lesson we will

- ✓ Identify the different sources of power consumption in embedded systems.
- ✓ Look at ways to measure power consumption.
- ✓ Study several different methods for managing power consumption.

Introduction

Today more and more embedded applications

Targeted towards

Small hand held or other types of portable devices

Common thread through all such applications

Need to have long battery life

Translates to low power consumption

Power consumption can be attacked in several ways

Certainly hardware solution

Low power devices

Turning portions of system off

ACPI – Advanced Configuration and Power Interface

Surprisingly have software contribution as well

Let's look at each and begin with a view into the software

Software

There are a number of places that we can attack

From software point of view

Initial places to look

- The algorithms that we use
- Location of code
 - Memory accesses can have significant impact on power
- Using software to control subsystems

As we have been stating

To analyze then control particular aspect of performance

Must be able to measure that aspect

Both before and after modification

Measuring Power Consumption

For the moment

Will assume goal is to reduce power consumed by processor

Processor can be stand alone or softcore on FPGA

To such an end

Measuring power consumption is two step process

1. Identify the portion of code to be analyzed

Typically this will be a loop

Doesn't need to be

Measure the current consumed by the processor

While the code is being exercised

2. Modify the loop such that the code comprising the loop is disabled

Ensure that the compiler hasn't optimized loop out

Measure the current consumed by the processor

Once we have identified power consumed

Next step is to try to reduce if appropriate

Studies have identified several software factors

That contribute to processor power consumption

Among the contributors we find

- The kind of instruction
- The collection or sequence of instructions executed
- The locations of the instructions and their operands

Memory system and transfers in and out

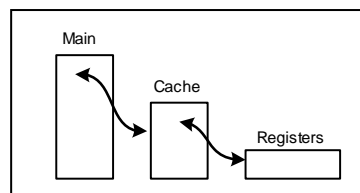
Have been shown to be most expensive operation

Performed by processor

Here memory is referring to main memory not cache

This is the DRAM in our system

Hierarchy will typically look like



Cache is high speed memory and is closer to the processor

Using simple addition operation as reference we find

<i>Operation</i>	<i>Relative Power Consumption</i>
16 Bit Add	1
16 Bit Multiply	3.6
8x128x16 SRAM Read - <i>Cache</i>	4.4
8x128x16 SRAM Write - <i>Cache</i>	9
I/O Access	10
16 Bit DRAM Memory Transfer - <i>Main</i>	33

Evident from table

Using cache can have significant affect on power consumption

Assumes a cache hit

Cache miss requires main memory access

SRAM generally consumes more power than DRAM

On per cell basis

Cache is generally SRAM

Want to optimize size of cache\

Want smallest that provides desired performance

Almost becomes empirical trade-off

Other optimizations

1. Power aware compilers

Take an instruction level view of problem

Modify schedule of bus activity

Bus drivers consume a lot of power

2. Use registers efficiently

Bring value into register and leave there for reuse

Don't move in and out

Read / write operations

3. Look for cache conflicts and eliminate if possible

Conflicts are data or instructions

That must be moved in or out of cache

For instruction conflicts

Rewrite if possible

May have to move code

Scalar data conflicts

- Move data to different locations
 - Arrayed data
 - Move to alternate location
 - Change access pattern

4. Unroll loops

- Must be careful that unrolled loop doesn't result in cache misses

5. Eliminate recursive procedures or repeated function calls where possible

- Eliminates overhead of function call

Hardware

Another technique for managing the power consumption

- In embedded applications

Draws from familiar schemes used at home

- Turn off the portions of system not being used

Such a scheme has been used for years

- In space program

- Satellites

- Orbital and interplanetary

- Shuttle

- Earlier Mercury, Gemini, Apollo

- Therein hardware

- Battery powered

- Must be recharged

- Done via solar panels of one form or another

Today can fly from Seattle to Japan in 14 hours

- Laptop computer or other such tools

- Typical battery life 3-5 hour

- Yep a laptop is still an embedded application

We are in continuous race between

- Battery technology

- Demand for more and more powerful features

- All such features require power

Power Management Schemes

To begin to address problem

As part of design

Formulate power management strategy

✓ On one extreme

Turn power off

In such state

Power consumption limited to leakage

As with other metrics

➤ Sets a lower bound on consumption

✓ Opposite extreme

Power to all parts of system on and all parts operating

In such state

Power consumption approaches maximum

Such condition sets upper bound

Softer than lower bound

See earlier discussion on software affects

Goal is somewhere in middle

Governed once again by requirements specification

Based upon such a goal

We segregate system components into two categories

Those that must remain powered up

Referred to as static components

Those that may be powered down

Referred to as dynamic components

Such a scheme sounds simple ...and is at the high level

Like everything else we are doing

We have certain trade-offs

We must

1. Decide which portions of the system to power down

These may be

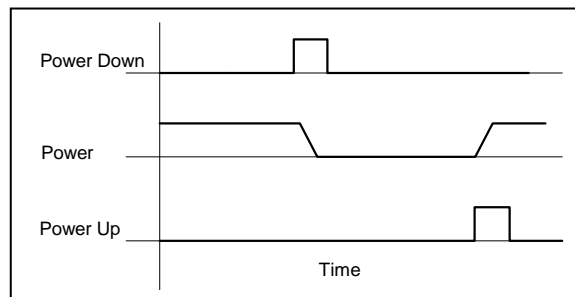
All dynamic components

Subgroups based up need or not need

2. Recognize that components cannot be shut down instantly

3. Recognize that components cannot be powered up instantly

These factors can be expressed with a simple first cut graphically as



Let's consider a topographic mapping satellite application

As satellite is circling the earth collecting data

Data is sent to ground station at known points in orbit

When over appropriate station

No reason to keep transmitter powered up

When not in position to transmit

Further timing of orbit known with sufficient resolution

Know in advance when will need to transmit

After passing ground station

Shut down transmitter

Re-enable shortly before reaching next download point

Locations of each ground station known in advance

1. Such fixed schedule scheme is among simplest

Can be very effective

Observe similar to

Round robin schedule with no preemption

2. Next level of sophistication

Recognize that schedule may not be fixed

Problem now moves from deterministic to probabilistic

Use knowledge of

Current history

Understanding of problem

To anticipate when to shut dynamic portions of system down

Denoted predictive shutdown

Observe that such a scheme commonly used in

Branch prediction logic in instruction prefetch pipeline

Managing if-else constructs
Using such a scheme
Subject to shutdown or restart too early

3. Related idea

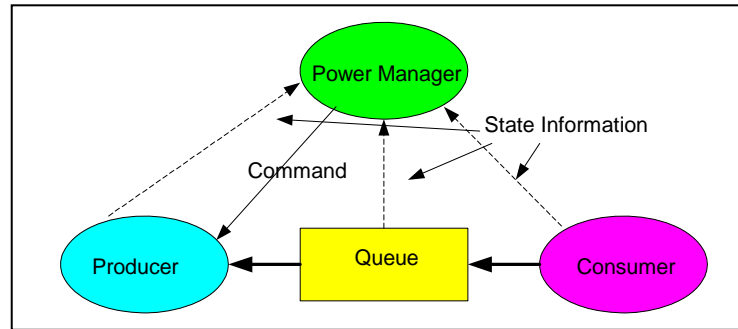
Rather than set schedule
Control algorithm with associated timer
Monitors activities of devices to be dynamically controlled
If timer expires
Device is powered down
Device reactivated on demand
We've already used such a scheme in a watchdog timer

4. Next level of sophistication

Draws from basic queuing theory
Under such a scheme we have
A resource or producer
Service provided by system whose power is being controlled
A consumer
Portion of the system that needs the service
A queue of service requests
A power manager
Monitors behaviour of system
Producer
Consumer
Queue
Can build schedule using Markov modeling
Statistical technique
Forecast future behaviour of variable or system
Who's current state of behaviour
Does not depend upon past
Random like flipping an honest coin

Approach maximizes
System computational performance
While satisfying specified power budget

Simple data / control flow diagram appears as



Let's look at an example of simple power management

- ✓ Queue
 - Queue of requests for power
- ✓ Consumer
 - Issues requests for power

Operating system responsible for dynamically controlling power

In simple I/O subsystem

Dynamically controlled portion

Supports two modes

Off and On

Dynamic subcomponents

Terminology

joule = 1 watt sec

power = joule / sec

Consume 10 watts when on and 0 watts when off

10 watts when on then in 25 sec since always on

$10 \text{ watts} \cdot 25 \text{ sec} = 250 \text{ joules}$

Switching

2 seconds and 40 joules

Switch from OFF state to ON state

1 second and 10 joules

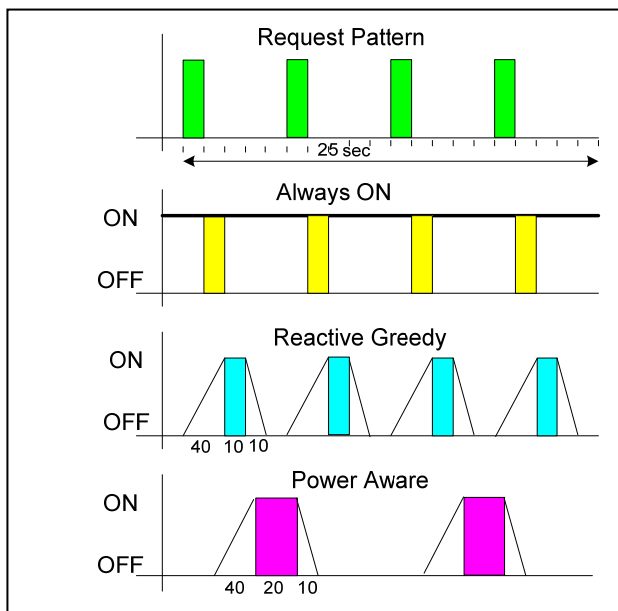
Switch from ON to OFF

Request has period of 25 seconds

Graphically we illustrate 3 alternative schemes

Observe that we have the same average throughput

Substantially reduced power consumption



Policy	Energy Consumed in 25 sec	Average Latency Per Request
Always ON	250 J (10 x 25 sec) (10 when on x 25)	1 sec
Reactive Greedy	240 J (40+10+10)x2 (60x4)	3 sec (2+1)x4/4
Power Aware	140 J (70 x 2)	2.5 sec (2+1)/2 + 1

5. Advanced Configuration and Power Interface – ACPI

Industry standard power management

Initial application was to PC

More specifically windows

Currently targeted to wider variety of operating systems

Standard

Provides some basic power management facilities

Provides interface to the hardware

Software more specifically the operating system on a system

Provides power management module

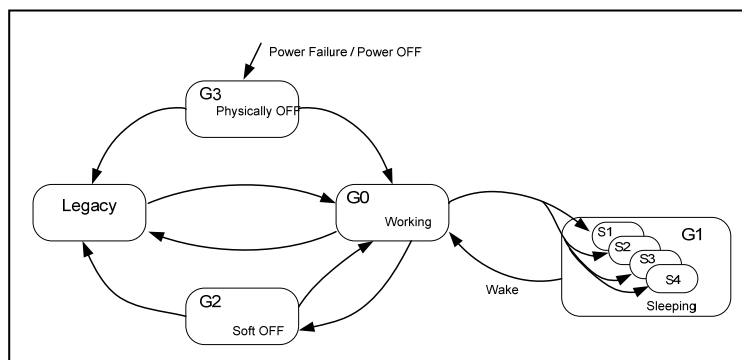
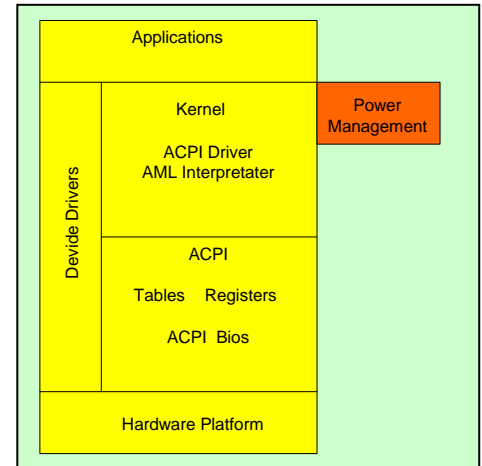
It is the responsibility of the OS

Specifying the power management policy for the system

The operating system uses the ACPI module
 To send required controls to the hardware
 Monitor the state of the hardware
 As input to the power manager

We express the scheme in the following

- Block diagram
- State diagram



Standard supports five global power states

1. G3 – hard off or full off
 Defined as physically off state
 System consumes no power
2. G2 – soft off requires full OS reboot to restore system to full operational condition

Substates

- S1 – low wake-up latency
 Ensures no loss of system context
- S2 – low wake-up latency state
 Has loss of CPU and system cache state
- S3 – low wake-up latency state
 All system state except for main memory is lost
- S4 – lowest power sleeping state
 All devices are off

3. G1 – sleeping state
 System appears to be off
 Time required to return to working condition
 Inversely proportional to power consumption

4. G0 – working state in which system is fully usable
5. Legacy state
 - a. System does not comply with ACPI

Trade-offs

Often times performance is optimization issue

Involves trading several contradictory requirements

Speed

Memory size

Cost

Weight

Power

Time must be spent up front to thoroughly understand

Application

Constraints

Summary

In this lesson we

- ✓ Identify the different sources of power consumption in embedded systems.
- ✓ Look at ways to measure power consumption.
- ✓ Study several different methods for managing power consumption.